

Daksh

dakshwadhwa2007@gmail.com • +91-9728928066 • [LinkedIn](#) • [GitHub: Dakshhhhh-ops](#) • [LeetCode](#)

Education

National Institute of Technology Hamirpur

B.Tech in Electrical Engineering

Aug 2024 – June 2028

CGPA: **8.78/10** | Dept. Rank: **13 / 125**

Technical Skills

Languages	Python, C++
AI / ML	NumPy, Pandas, Scikit-learn, TensorFlow, Keras, NLP
Generative AI	LangChain, RAG Pipelines, Groq API, Google Gemini API, Prompt Engineering, HuggingFace Embeddings, FAISS, Vector Stores
Backend/Deploy	FastAPI, Flask, Streamlit, REST APIs, Docker
Core CS	DBMS, Operating Systems
Tools	Git, GitHub, MLflow, Jupyter Notebook, VS Code, Figma

Projects

SpamShield-MLOps SMS Spam Classifier

[\[GitHub\]](#) — [\[Live Demo\]](#) — [\[API\]](#)

Python · Scikit-learn · FastAPI · Streamlit · MLflow · NLTK · TF-IDF · Docker

- Built a production-ready classifier with **decoupled architecture**: FastAPI backend on Render and Streamlit frontend communicating via REST API with real-time confidence scores.
- End-to-end NLP pipeline: lowercasing → tokenization → stopword removal → Porter stemming → TF-IDF → Naive Bayes; trained on UCI SMS Spam Collection (5,574 messages).
- Tracked 5+ MLflow experiments; best model achieved **97.5% accuracy and 100% precision**; containerized with Docker for reproducible deployment.

Financial Document RAG

[\[GitHub\]](#)

Python · LangChain · Groq LLM · FAISS · HuggingFace Embeddings · FastAPI · Streamlit

- Developed an end-to-end Retrieval-Augmented Generation (RAG) application for financial documents using LangChain, FAISS, Hugging Face, Groq, FastAPI, and Streamlit.
- Designed scalable indexing and retrieval workflows with recursive and table-aware chunking, semantic embeddings, persistent vector databases, and hybrid lexical-semantic search.
- Implemented citation-aware responses, retrieval evaluation metrics, and a production-ready backend with modular architecture.

MLOps Student Performance Predictor

[\[GitHub\]](#)

Python · Scikit-learn · Flask · AWS Elastic Beanstalk · AWS CodePipeline · Docker

- Architected a **modular ML pipeline** with separated data ingestion, transformation, and training components; deployed on AWS Elastic Beanstalk with CI/CD via AWS CodePipeline.
- Trained and compared 8+ models (Linear Regression, Ridge, CatBoost, XGBoost); best model achieved **0.897 R^2** with custom preprocessing for categorical encoding and feature scaling.
- Structured codebase with custom logging, exception handling, and artifact persistence; every push to main triggers automated redeployment.

Laptop Price Predictor

[\[GitHub\]](#) — [\[Live Demo\]](#)

Python · Pandas · Scikit-learn · Random Forest · Streamlit · Feature Engineering

- Engineered features from 1,300+ listings (PPI from resolution, CPU/GPU brand extraction, OS grouping, log-transformed target); trained Random Forest with cross-validated tuning.
- Achieved **R^2 of 0.89 and MAE of Rs.8,500** on test set; deployed live Streamlit app for instant INR price prediction across 13 hardware specs with feature importance visualization.

Achievements

- Winner** — Electron 8.0 Hackathon (IQAI Track); awarded **\$400 cash prize** for best AI/ML solution among 120+ competing teams.
- Smart India Hackathon (SIH)** — Selected among **College Top 40 Teams** out of 200+ college submissions.
- Solved **165+ problems** on LeetCode spanning data structures, algorithms, and dynamic programming.

Positions of Responsibility

Executive — Team .EXE, NIT Hamirpur

Jan 2025 – Present

- Led execution of technical projects and coordinated cross-functional student teams across development, design, and operations; mentored 30+ participants at large-scale campus tech events.

Executive — GDG on Campus, NIT Hamirpur

Aug 2024 – Present

- Organized AI/ML workshops and developer sessions for 50+ students covering ML fundamentals, model deployment, and cloud tools; coordinated community-driven technical projects.